Job Description: AI Engineer

Position Overview

We are looking for an analytical, execution-oriented AI Engineer to design, implement, and maintain advanced machine learning models and large language model architectures within our production platforms. This position sits at the intersection of core data science and software engineering, requiring the implementation of production-grade fine-tuning workflows, the optimization of complex data engineering pipelines, and the deployment of low-latency model inference systems.

Core Roles & Responsibilities

- Design, test, and productionize machine learning algorithms, deep learning models, and complex semantic search frameworks.
- Build and maintain optimized infrastructure for hosting Large Language Models (LLMs), integrating Retrieval-Augmented Generation (RAG) loops.
- Develop high-throughput, automated data pre-processing and pipeline systems to feed machine learning model pipelines.
- Optimize inference runtime parameters to ensure models hit strict latency and memory footprint constraints in production environments.
- Deploy, instrument, and track system health metrics for live AI models using modern, automated MLOps methodologies.

Required Qualifications & Experience

- Bachelor's or Master's degree in Computer Science, Data Science, Artificial Intelligence, or a strongly quantitative mathematical field.
- 2+ years of production experience training, deploying, or fine-tuning machine learning models in a cloud enterprise environment.
- Advanced command of Python and its core deep learning development toolkits, including PyTorch, TensorFlow, and Hugging Face Ecosystem.
- Hands-on expertise interacting with vector databases (such as Pinecone, Milvus, Qdrant, or Chroma) and orchestration tooling (LangChain, LlamaIndex).
- Solid software engineering fundamentals, including version control (Git), Docker containers, and RESTful/gRPC API delivery.

Preferred & Nice-to-Have Skills

- Familiarity with distributed machine learning training architectures and compute cluster orchestration using Ray, Spark, or Slurm.
- Experience deploying automated AI monitoring and orchestration workflows with MLflow, Kubeflow, or Weights & Biases.

- Contributions to open-source machine learning projects or published academic research in AI domains.